

Night Owl

Oliver Katz

June 8, 2013

This tutorial assumes that you have a basic knowledge of assembly. I'd recommend taking a look at a NASM tutorial or something if you want to know more. I hope this helps you with learning Night Owl. Thanks!

1 Design

Night Owl is a virtual machine which uses just-in-time (JIT) compilation to generate native machine code at runtime. Night Owl's kernel, the thing that actually generates the machine code, can be accessed through assembly easily as it is written in pure C and is accessible as a shared object (dynamically-linked library), allowing Night Owl programs to generate other Night Owl programs within themselves with a very low overhead.

The kernel contains little more than the code necessary for emitting machine code instructions into a memory block. The extended C++ library (contained in `build.h`) provides utilities for easier generation of code as well as optimizations for that code. This is done in a single pass (optimizations included).

Night Owl is designed to be flexible, in that it can run at many levels of abstraction and overhead; it can be as minimal as you want while still retaining some basic high-level features.

2 Using the Kernel

Night Owl is a register machine, which means that you need to use registers in your generated code. You can do this automatically with the `Builder` class, but not in the kernel. The first thing you need to do, is include the necessary headers.

```
#include <nightowl/nightowl.h>
```

To use the kernel, you need to create an `o_jit` structure, which contains the generated machine code.

```
o_jit jit = o_jit_init();
```

There are a group of functions called emitters. They all have the prototype `void emitter(o_jit *jit, unsigned char r0, unsigned char r1, int a);`, where `a0` is the first register parameter, `r1` is the second, and `a` is the integer parameter. If the emitter doesn't use a particular parameter, then it will still require you to provide it, but its value will not be taken into account.

2.1 Emitters

The first emitter, for example, is for the NOP instruction. Simply put, it does nothing. It simply skips to the next CPU cycle. You can emit this instruction with:

```
o_emit_nop(&jit, 0, 0, 0);
```

Actually, none of the parameters are used, so the following emittance would be equivalent:

```
o_emit_nop(&jit, 1, 2, 3);
```

The following is a reference of all of the instructions, relative to standard assembly instructions:

Emitter	Assembly Instructions	Description
o_emit_nop	???	Skips a CPU cycle and does nothing
o_emit_const	mov, movi	Store a into r0
o_emit_mov	mov	Move the value of r1 into r0
o_emit_mov_disp	mov, movd	Dereference the address in (r1 + a)
o_emit_ld	mov, ld	Dereference the address in r1 and store into r0
o_emit_st	mov, st	Store the value in r1 to the address in r0
o_emit_add	add	Add r0 to r1 and store in r0
o_emit_sub	sub	Subtract r1 from r0 and store in r0
o_emit_mul	mul	Multiply r0 by r1 and store in r0
o_emit_and	and	Logical-AND r0 and r1 and store in r0
o_emit_or	or	Logical-OR r0 and r1 and store in r0
o_emit_xor	xor	Logical-EXCLUSIVE-OR r0 and r1 and store in r0
o_emit_pusharg	push, pusha, pusharg	Push the value of r0 to the argument stack
o_emit_poparg	pop, popa, poparg	Pop the top value of the argument stack into r0
o_emit_leave	leave	Leave a function (must be put before the return instruction)
o_emit_ret	ret	Return the value of the EAX register (register 0) from the current function
o_emit_int	int	Call a system interrupt (relies on the values of EAX and EBX)
o_emit_cmp	cmp	Compare r0 and a for use in the branch instructions (like breq)
o_emit_breq	je	If the last values compared are equal, jump to the code position a
o_emit_brne	jne	If the last values compared are unequal, jump to the code position a
o_emit_brlt	jl	If the last register compared is less than the integer it is compared to, jump to the code position a
o_emit_brle	jle	If the last register compared is less than or equal to the integer it is compared to, jump to the code position a
o_emit_brgt	jg	If the last register compared is greater than the integer it is compared to, jump to the code position a
o_emit_brge	jge	If the last register compared is greater than or equal to the integer it is compared to, jump to the code position a
o_emit_jmp	jmp, jump	Jump to the code position a
o_emit_call	call	Call the function pointer a
o_emit_fld	fld	Push the double at address r0 to the floating-point stack
o_emit_fst	fst	Pop the double at the top of the floating-point stack to the address r0
o_emit_fadd	fadd	Add the top two doubles on the floating-point stack
o_emit_fsub	fsub	Subtract the top two doubles on the floating-point stack
o_emit_fmul	fmul	Multiply the top two doubles on the floating-point stack
o_emit_fdiv	fdiv	Divide the top two doubles on the floating-point stack
o_emit_fcmp	???	Compare the top two doubles on the floating-point stack (standard branch instructions still work with this)
o_emit_fchs	fchs	Change the sign of the top double on the floating-point stack
o_emit_fst_rot	fdecfpst	Pop the top element of the floating-point stack and push it to the bottom
o_emit_fnop	???	Floating point operations run on an alternate sub-processor called the FPU - this instruction is skips a cycle on that sub-processor and does nothing
o_emit_ftan	???	Take the tangent of the top double on the floating-point stack and push it to the top
o_emit_fatan	???	Take the arc-tangent of the top double on the floating-point stack and push it to the top
o_emit_fsin	???	Take the sine of the top double on the floating-point stack and push it to the top
o_emit_fcos	???	Take the cosine of the top double on the floating-point stack and push it to the top
o_emit_fsincos	???	Take the sine and cosine of the top double on the floating-point stack and push both to the top (faster than calling both instructions individually)
o_emit_fsqrt	???	Take the square root of the top double on the floating-point stack and push it to the top

2.2 Registers

There are eight registers that you can use: EAX, ECX, EDX, EBX, ESP, EBP, ESI, and EDI. Most of these serve special purposes some of the time, while some of them are general purpose. EAX, ECX, EDX, and EBX are generally quite general purpose, except EAX stores the information generated by comparison instructions (cmp and fcmp), ret returns the value stored in EAX, and int uses EAX and EBX. ESP points to the top of the stack at all times. EBP is the base register for functions and argument handling. ESI and EDI are not used within Night Owl, and cannot be used for general purpose operations. You can access these registers with the following symbols:

```
O_EAX O_ECX O_EDX O_EBX O_ESP O_EBP O_ESI O EDI
```

2.3 Examples

If you want to make a C-compatible function that returns a sum of its two integer arguments, you can do it as follows:

```
o_jit jit = o_jit_init();
o_emit_pusharg(&jit, O_EBP, 0, 0);
o_emit_mov(&jit, O_EBP, O_ESP, 0);
o_emit_mov_disp(&jit, O_EAX, O_EBP, 8);
o_emit_mov_disp(&jit, O_EBX, O_EBP, 12);
o_emit_add(&jit, O_EAX, O_EBX, 0);
o_emit_leave(&jit, 0, 0, 0);
o_emit_ret(&jit, 0, 0, 0);
```

If you want to make an if-statement, you can do that as follows:

```
// if 5 == 5,
o_emit_const(&jit, O_EBX, 0, 5);
o_emit_cmp(&jit, O_EBX, 0, 5);
o_emit_brne(&jit, 0, 0, 5); // branch to else-segment

// then-segment:
o_emit_const(&jit, O_ECX, 0, 1);
o_emit_jump(&jit, 0, 0, 6);

// else-segment: (code position: [5])
o_emit_const(&jit, O_ECX, 0, 0);

// end of if-statement (code position: [6])
// ecx should contain 1, because the if statement went into its then-segment
```

After we are done emitting our code, we need to run it. We can do this using a function pointer to the emitted code block in memory:

```
void *func = o_jit_fp(&jit);
```

This function pointer needs to be cast into the proper function pointer type. If you are returning something from your function, you will want a return value. If the returned data is larger than the return type, a segfault should occur. If you pass the wrong number of arguments, a segfault should occur. Here's an example of how to call the first example seen above:

```
int (*func)(int, int) = (int (*)(int, int))o_jit_fp(&jit);
int rtn = func(5, 2);
// rtn is the return value of the function, and will be 7
```

Or if we have defined multiple functions, you need to add the code position of the beginning of the function to the function pointer. To get the code position at the head of the generation at the current moment, use `o_jit_label(&jit);`.

```
o_jit jit = o_jit_init();
// generate some stuff
unsigned int func_label = o_jit_label(&jit);
// generate the function we want to call
void *func = o_jit_fp(&jit)+func_label;
// call func
```

3 The Builder Class

The Builder class is defined within `builder.h`. It is written in C++ and contains a higher overhead than the core kernel, but it can automate some common tasks and provide optimizations. The most important optimizations it provides are function inlining and lazily-evaluated register and floating-point operations.

Here is an example of how lazily-evaluated operations work. If you have two registers, EAX and EBX and you add their contents together, you get a result which takes three CPU cycles to calculate. If you know the values of EAX and EBX within the program (they are constant), you don't need to run the calculation at run time, you can run it at compile time by finding their sum and replacing the three instructions with one: store the integral result of the sum in EAX as a constant.

We first need to include the correct header:

```
#include <nightowl/builder.h>
```

Everything in this header is contained within the `nightowl` namespace. The following examples just assume that it is used, right out of the box.

Let us create a Builder class:

```
Builder b = Builder();
```

We can emit instructions directly into the builder:

```
b.emit(_add, 0_EAX, 0_EBX, 0);
```

We can also store them temporarily in `BuilderInstructions` (this has the same result):

```
BuilderInstruction inst = BuilderInstruction(_add, 0_EAX, 0_EBX, 0);
b.emit(inst);
```

And lastly, we can store many instructions as `BuilderInstructionPages`:

```
BuilderInstructionPage page = BuilderInstructionPage();
page << BuilderInstruction(_add, 0_EAX, 0_EBX, 0);
b.emit(page);
```

If we store multiple functions in a `BuilderInstructionPage`, as long as we use `function("theFunctionName");` at the beginning of each function, we can get the function segment with `page.getFunctionSegment("theFunctionName");`.

We can use Builders to generate machine code using the kernel, or we can output the generated instructions to a `BuilderInstructionPage` with:

```
b.outputToPage(); // enable page output
// use builder to generate code
BuilderInstructionPage page = b.page();
```

3.1 Function Inlining

As long as page output (see directly above) is enabled and we have only defined a single function within the builder, we can convert the contents of the builder to an instruction page as an inline function call. You can specify the arguments as integer references:

```
// generate function here
```

```
int arg0 = 5;
int arg1 = 10;
```

```
vector<int *> tmp;
tmp.push_back(&arg0);
tmp.push_back(&arg1);
```

```
BuilderInstructionPage inlinedCode = b.convertToInline(tmp);
```

You'll probably have to end up creating a new builder class to accomplish this. If output to page is enabled in this class, the extra overhead will be minimal.

3.2 Function-Generation Utilities

Instead of generating all of the instructions for generating function argument handling in a C-compatible way (as shown above in the kernel examples in section 2.3), we can use the following:

```
b.emitFunction();
unsigned char r0 = b.emitFunctionArgument(4); // 4 is the size of the argument in bytes (this is an int)
unsigned char r1 = b.emitFunctionArgument(4);
b.emit(_add, r0, r1, 0);
b.emit(_mov, 0_EAX, r0, 0); // r0 is probably eax, but to make sure (if it is, this instruction won't be g
b.emitFunctionReturn();
```

3.3 Relocation Utilities

Jumps and branches often jump to code positions that aren't generated yet. We can use the relocation utilities to handle this automatically.

```
b.emit(_const, 0_EBX, 0, 5);
b.emit(_cmp, 0_EBX, 0, 5);

b.emitRelocatableBranch("ifStatement", "elseSegment");
b.emit(_const, 0_ECX, 0, 1);
b.emitRelocatableJump("elseJump", "endIfStatement");

b.emitRelocatableLabel("elseSegment");
b.emit(_const, 0_ECX, 0, 0);

b.emitRelocatableLabel("endIfStatement");
```

Now, when execution reaches the branch "ifStatement", it will jump forwards to "elseSegment" in the code.

3.4 Lazily-Evaluated Operations

Let us start with the difference between dirty and pure registers. A pure register is one to which we know the value at compile-time. A dirty one is any other register.

If we know the value of a register at compile time, then we can operate on it at compile time, instead of run time. This is a simple optimization. You would normally see it in compilers operation on abstract syntax trees. Since we are running this same optimization on assembly code at the same time as it is being generated, it takes the form of pure and dirty registers.

Here is an example of a pure register with the value of 5:

```
unsigned char r = b.regAlloc(0_PURE);  
b.regValue(r, 5);
```

This register IS NOT emitted into the machine code yet. To do that, we need to call `b.emitReg(r)`; This will make `r` dirty. We can operate on pure registers with:

```
b.emitOperation(_add, r0, r1); // assuming we have registers r0 and r1
```

`emitOperation` works on both pure and dirty registers. We can push pure or dirty registers to the function argument stack with:

```
b.emitPushArg(r0);
```

We can do the same thing with floating-point stack elements.

```
b.fpDirty(0, 0_PURE); // make the top of the floating-point stack pure  
b.fpValue(0, 5.0); // set the value  
b.emitFp();  
b.emitFloatingOperation(_fchs);
```

You will notice that there is no example of adding new floating-point stack elements. You can do this with the constant and reference utilities seen directly below.

3.5 Constant and Reference Utilities

We can add constant integers (32-bits) with:

```
unsigned char r = b.emitConstI32(5);
```

Or a reference to an integer with:

```
unsigned char r = b.emitConstI32(&anInteger);
```

We can pull the value back to a reference with:

```
b.emitPullI32(r, &anInteger);
```

We can do the same things with the floating-point stack:

```
b.emitRefF64(&aDouble);  
b.emitPullF64();
```

3.6 Error Handling

If, while emitting machine code, errors occur, you can find them with the following functions:

```
while(b.countErrors() > 0)
{
    Builder::error e = b.popError();
}
```

e.c is the errorcode. **reg0AboveLimit** or **reg1AboveLimit** means that one of the register arguments is referencing a non-existent register, for example there is no register 57. It is recommended that only **0_EAX**, **0_ECX**, ... (see list of registers above) are used for safety. **nopWarning** is used whenever a NOP instruction is used, since that instruction wastes a CPU cycle. **nullDisplacement** is used when an instruction that should jump execution to another place in the code jumps to itself. This means infinite loops (note: just because you don't find this error code after generation does not mean that there are no infinite loops in your code as it does not check for mutually recursive jumps). **operatingDuplicateRegisters** that requires two different registers only gets two references to the same register. **nullCall** is when you try to call a null function pointer. These error codes will not stop generation and are sometimes not fatal. They are, however, important.

4 Conclusion

I hope this tutorial has helped you learn how to use Night Owl. Please look at the documentation for more information on use or email me any questions that you may have at: olivetti.katz@gmail.com. Thank you for reading.